

BostonMarathonChallenge_Project

September 2, 2019

```
[404]: url = 'https://www.xxx'
```

0.0.1 Unsupervised Learning Challenge Build your own NLP model

For this challenge, you will need to choose a corpus of data from nltk or another source that includes categories you can predict and create an analysis pipeline that includes the following steps:

1. Data cleaning / processing / language parsing
2. Create features using two different NLP methods: For example, BoW vs tf-idf.
3. Use the features to fit supervised learning models for each feature set to predict the category outcomes.
4. Assess your models using cross-validation and determine whether one model performed better.
5. Pick one of the models and try to increase accuracy by at least 5 percentage points.

Write up your report in a Jupyter notebook. Be sure to explicitly justify the choices you make throughout, and submit it below.

0.0.2 Questions that need answering:

1. What question are you trying to solve (or prove wrong) ?
Clustering for the Boston Marathon 2013 dataset.
2. What kind of data do you have? -> describe the source.. __Race results data, de-identified data from the race including:
 1. Age
 2. Gender
 3. Race times at intervals by 5k: 5k, 10k, 15k, 20k, half, 25k, 30k, 35k, 40k, and overall(finishing time).
 4. Division place
 5. Overall finishing place__
3. Do some EDA, plots
4. What's missing from the data and how do you deal with it?
Nothing is really missing from the dataset.

5. How can you add, change, or remove features to get more out of your data?

I added a new feature, called sustainer. A sustainer is a runner that maintains their xk mile pace standard deviation at or below 0.02. Basically at whatever speed they run, if they can be in the top 2% lowest standard deviation of runners, then they are a sustainer.

0.0.3 Test Report

Stuff 1 blah blah blah.

Stuff 2 Blah blah blah.

Stuff 3: Blah blah blah again

Key Learning. Blah blah blah again

```
[405]: # Constants
max_iterations      = 10                # set it to > 0 for determining the
      ↳ features importance
random_state        = 57
rows_in_training_set = 10000
rows_in_test_set    = 200000
test_size           = 0.10
train_size           = 0.90
rfc_test_size        = 50000
rfc_train_size       = 5000
sample_size         = 10
run_CountVectorizer  = False
run_TfidfVectorizer  = True
BegTimeStampNewlines = 3
EndTimeStampNewlines = 3
EndTimeStamp         = '\n'*EndTimeStampNewlines+'End'
BegTimeStamp         = 'End'+'\n'*BegTimeStampNewlines
SustainerSTDDEVLimit = 0.020

num_clusters = 3
target_column = 'overall_scaled'
xcolumnname = 'racestd'
```

```
[406]: # Controls
flag_to_run_rf = False
flag_to_plot_them = False
flag_to_run_correlation_matrix = False
flag_to_run_features_importance = False
flag_to_run_gradient_boosting = False
flag_to_run_linear_regression = False
flag_to_run_logistic_regression = False
flag_to_run_lasso_regression = False
flag_to_run_ridge_regression = False
```

```

flag_to_run_svc = False
flag_to_run_vectorizer_nb = False
flag_to_run_sentiment_analyzer = False
flag_to_run_affinity_propagation = True
flag_to_run_kmeans = True
flag_to_run_mean_shift = True
flag_to_run_spectral_clustering = True
flag_to_run_elbow_plot = True

debug = False

```

```

[407]: import pandas as pd
import numpy as np
import scipy
import matplotlib.pyplot as plt

%matplotlib inline

import chardet
import datetime
from sklearn import datasets, ensemble, metrics, linear_model
from sklearn.utils import shuffle
from sklearn.naive_bayes import BernoulliNB, MultinomialNB
import time, sys
import seaborn as sns
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.svm import SVC
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn import linear_model
from sklearn.model_selection import cross_val_predict, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import MinMaxScaler, normalize
from IPython.display import Markdown, display
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.decomposition import PCA
from sklearn import metrics
from sklearn.metrics import pairwise_distances, mean_squared_error
from sklearn.cluster import AffinityPropagation, KMeans, MeanShift,
    ↳estimate_bandwidth, SpectralClustering
from scipy.spatial.distance import cdist

```

```

[408]: # add this to a dictionary
# Constants
max_iterations      = 10                # set it to > 0 for determining the
    ↳features importance
random_state        = 57
test_size            = 0.10

```

```

train_size                = 0.90

begin_string = '\n'*3+'Begin'
end_string = 'End'+'\n'*3

# Regression/Classification control
Regression = False

print("Regression = {}".format(Regression))

```

Regression = False

```

[409]: # Display preferences.
%matplotlib inline
pd.options.display.float_format = '{:.3f}'.format
pd.set_option('display.max_rows', 1000)
pd.set_option('display.max_row', 1000)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)

[410]: def plot_time_to_complete():
    objects = ('BernoulliNB', 'MultinomialNB', 'Logistic Regression')
    y_pos = np.arange(len(objects))
    performance = [18,17,32]

    plt.bar(y_pos, performance, align='center', alpha=0.5)
    plt.xticks(y_pos, objects)
    plt.ylabel('Time in Minutes')
    plt.title('Yelp Sentiment Analysis Time to Complete')

    plt.show()

```

```

[411]: def file_stuff():

    #     global df
    path = "../..../.."
    filename = "Datafiles/bostonmarathon/results/2013/results.csv"
    print("fullfilename = {}".format(path+filename))
    df = pd.read_csv(path+filename)
    print("There are {} rows in this file.".format(df.shape[0]))
    return df

```

```

[412]: def dataset_cleanup(df):
    # data Cleanup

    #     global X, y

    df['gender_int'] = np.where(df['gender'] == 'M', 1, 0).astype(float)

```

```

    df['bib_int'] = df['bib'].replace(to_replace=r'[W|F]', value='-',
→regex=True).astype(int)
    kcolumns = ['5k', '10k', '20k', '25k', '30k', '35k', '40k', 'half']
    for kcol in kcolumns:
        df[kcol] = np.where(df[kcol] == '-', 0, df[kcol])
        df[kcol] = df[kcol].astype(float)
    df['5kpace'] = df['5k']/5.0
    df['10kpace'] = df['10k']/10.0
    df['20kpace'] = df['20k']/20.0
    df['halfpace'] = df['half']/21.095
    df['25kpace'] = df['25k']/25.0
    df['30kpace'] = df['30k']/30.0
    df['35kpace'] = df['35k']/35.0
    df['40kpace'] = df['40k']/40.0
    df['officialpace'] = df['official']/42.19
    # df['raceavg'] = ,axis=0).mean()
    df['racestd'] =
→df[['5kpace', '10kpace', '20kpace', 'halfpace', '25kpace', '30kpace', '35kpace', '40kpace', 'officialpace']]
→std(axis=1)
    df['raceavg'] =
→df[['5kpace', '10kpace', '20kpace', 'halfpace', '25kpace', '30kpace', '35kpace', '40kpace', 'officialpace']]
→mean(axis=1)
#     X = df[['age', 'gender_int', 'genderdiv', 'country',
→'official', 'racestd', 'raceavg']]
#     X = pd.get_dummies(X)

df.drop('ctz', axis=1, inplace=True)
df.drop('state', axis=1, inplace=True)
# these are the 2% sustainers. They can be running at any pace, but they
→are consistent!
    df['sustainer'] = np.where(df['racestd'] <= SustainerSTDDEVLimit, 1, 0).
→astype(float) # they sustained their pace very well for the race
    scaler = MinMaxScaler()

    scaler.fit(df[['age']])
    df['age_scaled'] = scaler.transform(df[['age']]).astype(float)

    scaler.fit(df[['overall']])
    df['overall_scaled'] = scaler.transform(df[['overall']]).astype(float)

    scaler.fit(df[['pace']])
    df['pace_scaled'] = scaler.transform(df[['pace']])

    scaler.fit(df[['official']])
    df['official_scaled'] = scaler.transform(df[['official']]).astype(float)

```

```

display('columns are now', df.columns)
#     df = fcn_MinMaxScaler(df, 'age', 'age_scaled')
#     df = fcn_MinMaxScaler(df, 'official', 'official_scaled')
X = df[['age_scaled', 'sustainer', 'gender_int', 'racestd']]
#     X = pd.get_dummies(X)

display("df columns cpt 92310: ", df.columns)

global target_column, xcolumnname, ycolumnname

#     target_column = 'overall_scaled'
#     xcolumnname = 'age_scaled'
ycolumnname = target_column

y = df[target_column]
printFormatted("target, y column is {}".format(target_column))

if debug == True:
    print_timestamp("X and y variables created")

printFormatted('we have cleaned up the dataframe.')
display_column_names('df values', df)
display_column_names('X values', X)
return df, X, y

```

```

[413]: def printFormatted(string):
        newline = '\n'
        display(Markdown(string))
        write_to_logfile(string+newline)

```

```

[414]: def fcn_MinMaxScaler(dataframe, orig_column, new_column):
        display("cp 1: In fcn_MinMaxScaler.  shape is:", dataframe.shape)
        scaler = MinMaxScaler()
        scaler.fit(dataframe[['{}'.format(orig_column)]])
        dataframe[['{}'.format(new_column)]] = scaler.transform(dataframe[['{}'.
→format(orig_column)]])
        display("cp 2: In fcn_MinMaxScaler.  shape is:", dataframe.shape)

        return dataframe

```

```

[415]: def plot_facet():
        g = sns.FacetGrid(data=df, col='stars')
        g.map(plt.hist, 'message_length', bins=50)

```

```

[416]: def write_to_logfile(message, mdformat=''):
        bufsize = 0
        with open('TestResults.md', 'a+') as the_file:
            the_file.write('{} {}'.format(mdformat, message))

```

```

[417]: def plot_model_accuracy():
    objects = ('BernoulliNB', 'MultinomialNB', 'Logistic Regression')
    y_pos = np.arange(len(objects))
    performance = [75.81, 85.98, 91.08]

    plt.bar(y_pos, performance, align='center', alpha=0.5)
    plt.xticks(y_pos, objects)
    plt.ylabel('Accuracy Percent')
    plt.title('Yelp Sentiment Analysis Accuracy')

    plt.show()

[418]: def print_timestamp(displaytext):
    import sys
    import datetime
    datetime_now = str(datetime.datetime.now())
    printFormatted("{:19.19}: In: {} {} ".format(datetime_now, sys._getframe(1).
    →f_code.co_name, displaytext))

[419]: def return_current_datetime():
    datetime_now = str(datetime.datetime.now())
    return datetime_now

[420]: # def printFormatted(string):
#     display(Markdown(string))

[421]: def data_demographics(dataframe, num_rows):

    display("dataframe.isnull().sum()", dataframe.isnull().sum())

    display("dataframe.columns\n", dataframe.columns)
    display("dataframe.head({})\n".format(num_rows), dataframe.head(num_rows))

    display("dataframe.sample({})\n".format(num_rows), dataframe.
    →sample(num_rows))
    display("dataframe.dtypes\n", dataframe.dtypes)
    display("dataframe.describe()\n", dataframe.describe())

[422]: def plot_them():
    for column in X_train.columns:
    #     plt.hist(X_train[column]*100, bins=40)
    plt.scatter(y_train, X_train[column]*100)
    plt.xlabel(column)
    plt.show()

[423]: def rfc_and_feature_importances(rf):    # Here we are using Random Forest
    →classifier method to determine the top 30 features.

    X_train, X_test, y_train, y_test = train_test_split(X, y,
    →test_size=test_size, train_size=train_size)

```

```

## Fit the model on your training data.
rf.fit(X_train, y_train)

## And score it on your testing data.
rf.score(X_test, y_test)

feature_importance = rf.feature_importances_

# Make importances relative to max importance.
feature_importance = 100.0 * (feature_importance / feature_importance.max())
sorted_idx = np.argsort(feature_importance)
cols=X.columns[sorted_idx].tolist()
cols=cols[::-1]
pos = np.arange(sorted_idx.shape[0]) + .5
plt.subplot(1, 2, 2)
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, X.columns[sorted_idx])
plt.xlabel('Relative Importance')
plt.title('Variable Importance')
plt.show()
#     print("We are returning these columns {}".format(cols))
return cols[:30] # return it sorted

```

```

[424]: def run_features_importance(rf,n):
# Here we will return the feature importances
    all_feature_important_columns = []

    for i in range(1,n):
        print_timestamp('running rfc iteration {} features importance for {}_
→times'.format(i,n))
        columns2 = rfc_and_feature_importances(rf)
#         columns2.extend('{}'.format(i))
        all_feature_important_columns = all_feature_important_columns + columns2
        #     print("all_feature_import_columns={}".
→format(all_feature_important_columns))

        print("\nBOD:\nall_feature_important_columns = {}\nEOD".
→format(sorted(all_feature_important_columns)))
        for feature in set(all_feature_important_columns):
            print_timestamp("the NOC of feature {} in all_feature_important_columns_
→is {}".format(feature, all_feature_important_columns.count(feature)))

```

```

[425]: def run_correlation_matrix():

    print_timestamp('Begin'+'\n'*3)

    # Setup the correlation matrix.

```



```

corrmat = X.corr()
print(corrmat)

# Set up the subplots
f, ax = plt.subplots(figsize=(12, 9))

# Let's draw the heatmap using seaborn.
sns.heatmap(corrmat, vmax=.6, square=True)
plt.show()

print_timestamp('\n'*3+'End')

```

```

[426]: def data_characteristics():

    printFormatted("#### Columns used in the dataset")
    display(df.columns)

    print("\n\n")
    printFormatted("#### Describe of the df dataset")
    display(df.describe())

    print("\n\n")
    printFormatted("#### Sample of 10 from the dataset")
    display(df.sample(sample_size))

    print("\n\n")
    printFormatted("#### Number of nulls in X")
    display(X.isnull().sum())
    print("\n\n\n")

```

```

[427]: def training_test_set(X, y):
    #     global X_train, X_test, y_train, y_test
    # Let's fit it with the RFC training set
    X_train, X_test, y_train, y_test = train_test_split(X, y,
→test_size=test_size, train_size=train_size, random_state=0)
    print("train_size = {}, X_train is {}, and y_train is {}".
→format(train_size, len(X_train), len(y_train)))
    print("test_size = {}, X_test is {}, and y_test is {}".format(test_size,
→len(X_test), len(y_test)))

    return X_train, X_test, y_train, y_test

```

```

[428]: def run_rf(rf):
    print_timestamp('Begin run_rf part 1')

    ## Fit the model on your training data.
    rf.fit(X_train, y_train)

```

```

## Let's score it with the training data set
train_score = rf.score(X_train, y_train)
print("Training score = {}".format(train_score))

## Let's score it with the test data set
test_score = rf.score(X_test, y_test)

print("Test score = {}".format(test_score))
#Let's run cross validate score with the training data set
#    cross_val_score(rf, X_train, y_train, cv=5)

print_timestamp('End run_rfr part 1')

```

```

[429]: def run_BernoulliNB(data, target):

    # Our data is binary / boolean, so we're importing the Bernoulli classifier.

    # Instantiate our model and store it in a new variable.
    bnb = BernoulliNB()

    # Fit our model to the data.
    bnb.fit(data, target)

    # Classify, storing the result in a new variable.
    y_pred = bnb.predict(data)

    # Display our results.
    print("Number of mislabeled points out of a total {} points : {}".format(
        data.shape[0],
        (target != y_pred).sum()
    ))

    ## Let's score it with the test data set
    test_score = bnb.score(data, target)

    print("Test score = {}".format(test_score))

```

```

[430]: def sentiment_analyzer(path, parameters, classifier, tfidf_parms):
    # path A = the old path
    # path B = the new path, no CountVectorizer at all

    # run block of code and catch warnings

    if debug == True:
        print_timestamp(BegTimeStamp+" running with path={}".format(path))

    global vectorized

```

```

vectorized = True

pipeline_array = []

if path == "A":
    if classifier == 'bnb':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', BernoulliNB(**parameters))
        ]))
    elif classifier == 'svc':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', SVC(kernel = 'linear', **parameters))
        ]))
    elif classifier == 'mlb':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', MultinomialNB(**parameters))
        ]))
    elif classifier == 'logit':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', LogisticRegression(**parameters))
        ]))
    elif classifier == 'rfc':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', ensemble.RandomForestClassifier(**parameters))
        ]))

elif path == "B":
    if classifier == 'bnb':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', BernoulliNB(**parameters))
        ]))
    elif classifier == 'svc':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', SVC(kernel = 'linear', **parameters))
        ]))
    elif classifier == 'mlb':
        pipeline_array.append(Pipeline([
            ('tfidf', TfidfVectorizer(**tfidf_parms)),
            ('clf', MultinomialNB(**parameters))
        ]))

```

```

elif classifier == 'logit':
    pipeline_array.append(Pipeline([
        ('tfidf', TfidfVectorizer()),
        ('clf', LogisticRegression(**parameters))
    ]))
elif classifier == 'rfc':
    pipeline_array.append(Pipeline([
        ('tfidf', TfidfVectorizer(**tfidf_parms)),
        ('clf', ensemble.RandomForestClassifier(**parameters))
    ]))

pipe = pipeline_array[0]

try:
    vect_name_list = str(pipe.named_steps['vect']).split('(')
    vect_name = "vect = {}", ".format(vect_name_list[0])
except:
    vect_name = ''

classifier_name_list=str(pipe.named_steps['clf']).split('(')
classifier_name=classifier_name_list[0]
tfidf_name_list = str(pipe.named_steps['tfidf']).split('(')
if len(tfidf_name_list) > 0:
    tfidf_name = tfidf_name_list[0]
else:
    tfidf_name = ''

printFormatted("#### Now running with: {} tfidf={} and clf={}
->{}\n\nparameters={} \n\n tfidf_parms={} ".format( vect_name,

->                                tfidf_name,

->                                classifier_name,

->                                return_current_datetime(),

->                                parameters,

->                                tfidf_parms

->                                ))
pipe.fit(X_train, y_train)

y_pred_class = pipe.predict(X_test)
y_pred_class2 = pipe.predict(X_train)

```

```

metrics_test_score = metrics.accuracy_score(y_test, y_pred_class)
metrics_train_score = metrics.accuracy_score(y_train, y_pred_class2)

printFormatted('### Metrics test accuracy score = {:.2%} with {}'.format(
    metrics_test_score, classifier_name))
printFormatted('### Metrics train accuracy score = {:.2%} with {}'.format(
    metrics_train_score, classifier_name))

if debug == True:
    printFormatted("Steps information: {}".format(pipe.steps))
    print_timestamp("Finished running pipeline with:\n{}: ".format(
        classifier_name))

return y_test, y_pred_class

print_timestamp(EndTimeStamp)

```

Let's try predicting with gradient boosting classification

```

[431]: def run_gradient_boosting():

    print_timestamp('Begin')

    clf = ensemble.GradientBoostingClassifier(**params)

    #Let's run cross validate score with the training data set
    cross_val_score(clf, X_train, y_train, cv=5)

    loss_function = 'deviance' # could be exponential
    depth_value = 8
    params = {'n_estimators': 500,
              'max_depth': 8,
              'loss_function': loss_function,
              'max_leaf_nodes': depth_value, # 8 worked best...
              'min_samples_leaf': depth_value * 3
              , 'random_state' : random_state
              }

    clf.fit(X_train, y_train)

    predict_train = clf.predict(X_train)
    predict_test = clf.predict(X_test)

    print_timestamp('End')

```

```

[432]: def run_svc():

    print_timestamp('\n'*3+'Begin run_svc')

```

```

# Let's do a linear Support Vector Classifier
print_timestamp('Running SVC(kernel=linear)')
svm = SVC(kernel = 'linear')

# Let's fit the training model
print_timestamp('Running svm.fit')
svm.fit(X_train, y_train)

# Let's score the training set
print_timestamp('Running svm.score for the training set')
svm.score(X_train, y_train)

# Let's score the test set
print_timestamp('Running svm.fit for the test set')
svm.score(X_test, y_test)

print_timestamp('\n'*3+'End run_svc')

```

```

[433]: def run_logistic_regression():
    print_timestamp('\n'*3+'Begin')

    lr = LogisticRegression(C=1e20, solver='lbfgs', max_iter=1000)

    print_timestamp('Running lr.fit for the training set')
    lr.fit(X_train, y_train)

    print_timestamp('Running lr.fit for the training set')
    print('\nR-squared simple model training set yields:')
    print(lr.score(X_train, y_train))
    print("here comes the test set")
    lr_score = lr.score(X_test, y_test)
    printFormatted("### Logistic Regression score={:.2%}".format(lr_score))

    print_timestamp('\n'*3+'End')

```

```

[434]: def run_linear_regression():

    print_timestamp('\n'*3+'Begin')

    regr = linear_model.LinearRegression()

    print_timestamp('Running regr.fit for the training set')
    regr.fit(X_train, y_train)

    print("\nCoefficients: \n", regr.coef_)
    print("\nIntercept: \n", regr.intercept_)
    print("\nR-squared for training data set:")

```

```

print(regr.score(X_train, y_train))

print("\nR-squared for test data set:")
print(regr.score(X_test, y_test))

print_timestamp('End run_linear_regression.\n\n')

print_timestamp('\n'*3+'End')

```

```

[435]: def run_ridge_regression():
    # Fitting a ridge regression model. Alpha is the regularization
    # parameter (usually called lambda). As alpha gets larger, parameter
    # shrinkage grows more pronounced. Note that by convention, the
    # intercept is not regularized. Since we standardized the data
    # earlier, the intercept should be equal to zero and can be dropped.
    print_timestamp('\n'*3+'Begin')

    ridgeregr = linear_model.Ridge(alpha=10, fit_intercept=False)
    ridgeregr.fit(X_train, y_train)
    print(ridgeregr.score(X_train, y_train))

    print_timestamp('\n'*3+'End')

```

```

[436]: def run_affinity_propagation(data, target):

    print_timestamp('\n'*3+'starting AffinityPropagation')

    print_timestamp('\n'*3+'Begin')

    ap = AffinityPropagation()
    #     ap = AffinityPropagation(damping=0.5,
    #                               max_iter=200,
    #                               convergence_iter=15,
    #                               copy=True,
    #                               preference=None,
    #                               affinity='euclidean',
    #                               verbose=False)

    model = ap.fit(data)
    pred = ap.predict(data)

    Z = merge_predict_and_cluster(data, target, pred) # let's merge the data,
    →dataframe, prediction, and the cluster

    # Pull the number of clusters and cluster assignments for each data point.
    cluster_centers_indices = ap.cluster_centers_indices_
    n_clusters_ = len(cluster_centers_indices)
    labels = ap.labels_

```

```

print('Estimated number of clusters: {}'.format(n_clusters_))

labels = model.labels_

print("from run_affinity_propagation {}".format(metrics.
→silhouette_score(data, labels, metric='euclidean'))))

print_timestamp('\n'*3+'finished with AffinityPropagation')

return Z, n_clusters_

```

```

[437]: def run_kmeans(data, target, K):

    print_timestamp('\n'*3+'Begin')
    print("running with number of clusters = {}".format(K))
    km = KMeans(n_clusters=K, random_state=42)

    #     pred = KMeans(n_clusters=K, random_state=42).fit_predict(data)
    pred = km.fit_predict(data)
    #     Z = pd.DataFrame()
    Z = merge_predict_and_cluster(data, target, pred) # let's merge the data,
→dataframe, prediction, and the cluster

    #     Z = pd.merge(data, pd.DataFrame(pred), left_index=True, right_index=True)
    #     display_column_names('first Z values', Z)
    #     Z.rename(columns={Z.columns[-1]: 'cluster'}, inplace=True)
    #     display_column_names('second Z values', Z)
    #     Z = pd.merge(Z, target, left_index=True, right_index=True)
    #     display_column_names('third Z values', Z)
    #     print("z columns are {}".format(Z.columns))

    if debug == True:
        print("the shape of Kmeans_pred is {}, and the shape of X is {}, and,
→the shape of Z is {}".format(pred.shape,

→                                     data.shape,

→                                     Z.shape))
        display(Z.head(100))
        display_column_names('Z below values', Z)

        count = Z.groupby(['cluster']).count()
        display("Z: Count by clusters are this:\n", count)

    return Z

    print_timestamp('\n'*3+'End')

```



```

[438]: def merge_predict_and_cluster(dataframe, target, predict):
        Z = pd.merge(dataframe, target, left_index=True, right_index=True)
        Z = pd.merge(Z, pd.DataFrame(predict), left_index=True, right_index=True)
        Z.rename(columns={Z.columns[-1]: 'cluster'}, inplace=True)

        return Z

[439]: def run_spectral_clustering(data, target, K):
        display_dataframe_shape('entering run_spectral_clustering, data has shape_
        of:', data)
        display_dataframe_shape('entering run_spectral_clustering, target has shape_
        of:', target)
        print_timestamp('\n'*3+'Begin')

        # for clusternum in range(2, K):
        print_timestamp("Running spectral_clustering with {} clusters.".format(K))
        n_clusters=K

        # Declare and fit the model.
        sc = SpectralClustering(n_clusters=K)
        sc.fit(data)

        #Predicted clusters.
        predict=sc.fit_predict(data)

        Z = merge_predict_and_cluster(data, target, predict) # let's merge the data_
        dataframe, prediction, and the cluster

        if debug == True:
            display_dataframe_shape('in run_spectral_clustering, Z has shape of:',
            Z)
            display_dataframe_shape('in run_spectral_clustering, target has shape_
            of:', target)
            display("the datatypes of Z are", Z.dtypes)

        # plt.scatter(Z['cluster'], Z[target_column], c=Z['cluster'])
        # plt.show()

        labels = sc.labels_
        print("from spectral clustering {}".format(metrics.silhouette_score(data,
        labels, metric='euclidean'))))

        # print('Comparing the assigned categories to the ones in the data:')
        # print(pd.crosstab(target,predict))

        print_timestamp('\n'*3+'End')

```

```
return Z
```

```
[440]: def do_the_elbow(X):
        printFormatted("## We are plotting the elbow method!")
        # calculate distortion for a range of number of cluster
        distortions = []
        for i in range(1, 11):
            km = KMeans(
                n_clusters=i, init='random',
                n_init=10, max_iter=300,
                tol=1e-04, random_state=0
            )
            km.fit(X)
            distortions.append(km.inertia_)

        # plot
        plt.plot(range(1, 11), distortions, marker='o')
        plt.xlabel('Number of clusters')
        plt.ylabel('Distortion')
        plt.show()

[441]: def plot_it_clusters(dataframe, xvalue, yvalue, title):

        if debug == True:
            display_dataframe_shape('entry received in plot_it_clusters', dataframe)
            display(dataframe.dtypes)

        data_demographics(dataframe, 10)

        plt.rcParams['figure.figsize'] = [xvalue, yvalue]
        plt.xlabel(xcolumnname)
        plt.ylabel(ycolumnname)

        df0 = dataframe[dataframe.cluster == 0]
        df1 = dataframe[dataframe.cluster == 1]
        df2 = dataframe[dataframe.cluster == 2]
        df3 = dataframe[dataframe.cluster == 3]
        df4 = dataframe[dataframe.cluster == 4]
        df5 = dataframe[dataframe.cluster == 5]

        plt.scatter(df0[xcolumnname], df0[ycolumnname], color='green')
        plt.scatter(df1[xcolumnname], df1[ycolumnname], color='red')
        plt.scatter(df2[xcolumnname], df2[ycolumnname], color='blue')
        plt.scatter(df3[xcolumnname], df3[ycolumnname], color='black')
        plt.scatter(df4[xcolumnname], df4[ycolumnname], color='magenta')
        plt.scatter(df5[xcolumnname], df5[ycolumnname], color='orange')
        plt.title(title)
        plt.show()
```

```

#     plt.scatter(km.cluster_centers_[0], km.cluster_centers_[
→, 1], color='purple', marker='*', label='centroid')

#     if type == 'KMeans':
#         plt.xlabel('Age')
#         plt.ylabel('Income ($)')
#         plt.legend()
#         plt.scatter(km.cluster_centers[:, 0],
#                     km.cluster_centers[:, 1],
#                     marker = '*',
#                     label = 'centroid')

```

```

[442]: def run_mean_shift(data, target):

    print_timestamp('\n'*3+'Begin')

    X_train = data

    # Here we set the bandwidth. This function automatically derives a
→bandwidth
    # number based on an inspection of the distances among points in the data.
    bandwidth = estimate_bandwidth(X_train, quantile=0.2, n_samples=500)

    # Declare and fit the model.
    ms = MeanShift(bandwidth=bandwidth, bin_seeding=True)
    if debug == True:
        display_dataframe_shape('this is the shape of data coming into
→run_mean_shift', data)
    ms.fit(data)

    if debug == True:
        display_dataframe_shape('this is the shape of target coming into
→run_mean_shift', target)
    pred = ms.predict(data)
    if debug == True:
        display_dataframe_shape('this is the shape of pred after predict in
→run_mean_shift', data)

    Z = merge_predict_and_cluster(data, target, pred) # let's merge the data
→dataframe, prediction, and the cluster

    # Extract cluster assignments for each data point.
    labels = ms.labels_

    print("from mean shift {}".format(metrics.silhouette_score(data, labels,
→metric='euclidean'))))

```

```

# Coordinates of the cluster centers.
cluster_centers = ms.cluster_centers_

# Count our clusters.
n_clusters_ = len(np.unique(labels))

print("Number of estimated clusters: {}".format(n_clusters_))

print_timestamp('\n'*3+'End')

return Z, n_clusters_

```

```

[443]: def vectorizer_nb(type_of_vectorizer):

    print_timestamp(BegTimeStamp)

    # 1. import and instantiate CountVectorizer (with the default parameters)

    # 2. instantiate CountVectorizer (vectorizer)

    #     X = df.message
    #     y = df.sentiment_label

    # split X and y into training and testing sets
    # by default, it splits 75% training and 25% test
    # random_state=1 for reproducibility

    #     X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1)

    # 3. fit & transform
    if type_of_vectorizer == 'Count':
        print("We are running with CountVectorizer")
        vectorizer = CountVectorizer()
        vectorizer.fit(X_train)
        vectorizer_method = 'CountVectorizer'
    elif type_of_vectorizer == 'Tfidf':
        print("We are running with TfidfVectorizer")
        vectorizer = TfidfVectorizer()
        vectorizer.fit_transform(X_train)
        vectorizer_method = 'TfidfVectorizer'

    # 4. transform training data
    X_train_dtm = vectorizer.transform(X_train)

    # equivalently: combine fit and transform into a single step
    # this is faster and what most people would do

```

```

X_train_dtm = vectorizer.fit_transform(X_train)

# 4. transform testing data (using fitted vocabulary) into a document-term
→matrix
X_test_dtm = vectorizer.transform(X_test)

# 1. import

# 2. instantiate a Multinomial Naive Bayes model
nb = MultinomialNB()

# 3. train the model
# using X_train_dtm (timing it with an IPython "magic command")

nb.fit(X_train_dtm, y_train)

# 4. make class predictions for X_test_dtm
y_pred_class = nb.predict(X_test_dtm)

# calculate accuracy of class predictions

met_test_score = metrics.accuracy_score(y_test, y_pred_class)
printFormatted('### With {} vectorizer, the metrics accuracy score = {:.
→2%}'.format(vectorizer_method,

→          met_test_score))

print_timestamp(EndTimeStamp)

```

```

[444]: def display_column_names(label, df):
        display("Label: {}: Column names are:".format(label), df.columns)

```

```

[445]: def display_dataframe_shape(label, df):
        display("Label: {}: Dataframe shape is:".format(label), df.shape)

```

```

[446]: def run_it(X_train, X_test, y_train, y_test, y):

#     file_stuff()

#     data_cleanup()

print_timestamp('\n'*3+'Begin')

if Regression == True:
    print_timestamp("We are running with a Regression model")
elif Regression == False:
    print_timestamp("We are running with a Classifier model")

```

```

else:
    print_timestamp("We have failed to set the Regression variable")
    sys.exit(main())

if flag_to_plot_them == True:
    plot_them()

if flag_to_run_features_importance == True:

    number_of_features_to_consider = 50
    params = {'n_estimators': 100}

    if Regression == True:
        print_timestamp('We are running RandomForestRegressor')
        rf = ensemble.RandomForestRegressor(**params)

    else:
        print_timestamp('We are running RandomForestClassifier')
        rf = ensemble.RandomForestClassifier(**params)

    run_features_importance(rf, number_of_features_to_consider)

if flag_to_run_correlation_matrix == True:
    run_correlation_matrix()

if flag_to_run_rf == True:
    #     params = {}
    params = {'n_estimators': 100}

    if Regression == True:
        rf = ensemble.RandomForestRegressor(**params)
        print_timestamp('We are running RandomForestRegressor')
    else:
        rf = ensemble.RandomForestClassifier(**params)
        print_timestamp('We are running RandomForestClassifier')

    run_rf(rf)

if flag_to_run_gradient_boosting == True:
    run_gradient_boosting()

if flag_to_run_linear_regression == True:
    run_linear_regression()

if flag_to_run_logistic_regression == True:
    run_logistic_regression()

```

```

if flag_to_run_svc == True:
    run_svc()

if flag_to_run_ridge_regression == True:
    run_ridge_regression()

if flag_to_run_vectorizer_nb == True:
    for vectorizer_iterator in ['Count', 'Tfidf']:
        vectorizer_nb(vectorizer_iterator)

if flag_to_run_kmeans == True:
    method = KMeans(
        n_clusters=num_clusters
#         ,random_state=42
#         ,init='random'
#         ,n_init=10
#         ,max_iter=300
#         ,tol=1e-04
    )
    df1 = run_kmeans(X_train, y_train, num_clusters)
    plot_it_clusters(df1, xvalue=16, yvalue=16, title="KMeans with number_
→of clusters = {}".format(num_clusters))
    display("next plot please")

if flag_to_run_affinity_propagation == True:
    display_column_names('columns of X_train going into_
→affinity_propagation: ', X_train)
    df2, ap_num_clusters = run_affinity_propagation(X_train, y_train)
    plot_it_clusters(df2, xvalue=16, yvalue=16, title="Affinity Propagation_
→with number of clusters = {}".format(ap_num_clusters))

if flag_to_run_mean_shift == True:
    df3, mean_shift_num_clusters = run_mean_shift(X_train, y_train)
    plot_it_clusters(df3, xvalue=16, yvalue=16, title="Mean Shift with_
→number of clusters = {}".format(mean_shift_num_clusters))

if flag_to_run_spectral_clustering == True:
    df4 = run_spectral_clustering(X_train, y_train, K=num_clusters)
    plot_it_clusters(df4, xvalue=16, yvalue=16, title="Spectral clustering_
→with number of clusters = {}".format(num_clusters) )

print_timestamp('End'+'\n'*3)

```

```
[447]: def main(entry_point):
```

```
    if entry_point == 0:
```

```

print_timestamp("Starting main()")
df = file_stuff()
data_demographics(df, 5)
display_column_names('post data_demographics of df', df)
df, X, y = dataset_cleanup(df)
display_column_names('post dataset_cleanup on X', X)
data_demographics(df, 5)
display_column_names('post data_demographics on X #2', X)
#     make_X_and_Y()
X_train, X_test, y_train, y_test = training_test_set(X, y)
display_column_names('after training_test_set: columns of X_train going_
→into affinity_propagation: ', X_train)
#     data_characteristics()
#     plot_time_to_complete()
#     plot_model_accuracy()
#     plot_facet()

if flag_to_run_elbow_plot == True:    do_the_elbow(X)
run_it(X_train, X_test, y_train, y_test, y)

print_timestamp("Ending main()")

```

[448]: main(0)

2019-07-25 23:22:55: In: main Starting main()

fullfilename = ../../../../Datafiles/bostonmarathon/results/2013/results.csv
There are 16164 rows in this file.

'dataframe.isnull().sum()'

25k	0
age	0
name	0
division	0
10k	0
gender	0
half	0
official	0
bib	0
ctz	15407
country	0
overall	0
pace	0
state	1463
30k	0
5k	0
genderdiv	0


```

20k          0
35k          0
city         1
40k          0
dtype: int64

```

```
'dataframe.columns\n'
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'ctz', 'country', 'overall'])
```

```
'dataframe.head(5)\n'
```

	25k	age	name	division	10k	gender	half	official	bib	ctz	country	overall
0	49.87	28	Cassidy, Josh R.	9	18.18	M	40.93	90.900	W1	NaN	CAN	10.00
1	77.27	30	Korir, Wesley	5	30.90	M	64.90	132.500	1	NaN	KEN	10.00
2	77.23	23	Desisa, Lelisa	1	30.90	M	64.92	130.370	2	NaN	ETH	10.00
3	50.50	32	Fearnley, Kurt H.	5	18.73	M	42.00	88.430	W2	NaN	AUS	10.00
4	48.75	39	Hokinoue, Kota	3	18.18	M	40.57	87.220	W3	NaN	JPN	10.00

```
'dataframe.sample(5)\n'
```

	25k	age	name	division	10k	gender	half	official	bib	ctz	country	overall
7162	121.72	34	Anderson, Rebecca J.	3348	47.58	F	101.25	235.550	8676	NaN	USA	10.00
133	85.67	27	Madut, Thomas G.	53	33.60	M	71.92	149.000	150	NaN	KEN	10.00
5523	109.72	55	Anderson, Nate	102	44.83	M	92.75	197.420	6697	NaN	USA	10.00
9782	122.00	38	St. Clair, Suzanne W.	2051	47.75	F	102.27	216.170	11870	NaN	USA	10.00
7437	121.45	56	Walls, Richard S.	557	47.68	M	102.15	223.350	9012	NaN	USA	10.00

```
'dataframe.dtypes\n'
```

```

25k          object
age          int64
name         object
division     int64
10k          object
gender       object
half         object
official     float64
bib          object
ctz          object
country      object

```

```

overall      int64
pace         float64
state        object
30k          object
5k           object
genderdiv    int64
20k          object
35k          object
city         object
40k          object
dtype: object

```

```
'dataframe.describe()\n'
```

	age	division	official	overall	pace	genderdiv
count	16164.000	16164.000	16164.000	16164.000	16164.000	16164.000
mean	41.638	1100.967	208.159	8429.373	7.947	4351.685
std	10.351	942.115	23.744	5052.024	0.906	2772.398
min	18.000	1.000	85.530	1.000	3.270	1.000
25%	34.000	363.000	191.727	4061.750	7.320	2032.750
50%	42.000	842.000	209.225	8247.500	7.980	4113.500
75%	49.000	1560.000	225.230	12662.250	8.600	6316.000
max	80.000	3834.000	284.230	17598.000	10.850	10648.000

```
'Label: post data_demographics of df: Column names are:'
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'ctz', 'c
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:334:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by MinMaxScaler.
```

```
    return self.partial_fit(X, y)
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:334:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by MinMaxScaler.
```

```
    return self.partial_fit(X, y)
```

```
'columns are now'
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'country'
```

```
'df columns cpt 92310: '
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'country']
```

target, y column is overall_scaled
we have cleaned up the dataframe.

```
'Label: df values: Column names are:'
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'country']
```

```
'Label: X values: Column names are:'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd'], dtype='object')
```

```
'Label: post dataset_cleanup on X: Column names are:'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd'], dtype='object')
```

```
'dataframe.isnull().sum()'
```

25k	0
age	0
name	0
division	0
10k	0
gender	0
half	0
official	0
bib	0
country	0
overall	0
pace	0
30k	0
5k	0
genderdiv	0
20k	0
35k	0
city	1
40k	0
gender_int	0
bib_int	0
5kpace	0

```

10kpace      0
20kpace      0
halfpace     0
25kpace      0
30kpace      0
35kpace      0
40kpace      0
officialpace  0
racestd      0
raceavg      0
sustainer    0
age_scaled   0
overall_scaled 0
pace_scaled  0
official_scaled 0
dtype: int64

```

```
'dataframe.columns\n'
```

```
Index(['25k', 'age', 'name', 'division', '10k', 'gender', 'half', 'official', 'bib', 'country', 'overall'])
```

```
'dataframe.head(5)\n'
```

	25k	age	name	division	10k	gender	half	official	bib	country	overall
0	49.870	28	Cassidy, Josh R.	9	18.180	M	40.930	90.900	W1	CAN	9
1	77.270	30	Korir, Wesley	5	30.900	M	64.900	132.500	1	KEN	5
2	77.230	23	Desisa, Lelisa	1	30.900	M	64.920	130.370	2	ETH	1
3	50.500	32	Fearnley, Kurt H.	5	18.730	M	42.000	88.430	W2	AUS	5
4	48.750	39	Hokinoue, Kota	3	18.180	M	40.570	87.220	W3	JPN	3

```
'dataframe.sample(5)\n'
```

	25k	age	name	division	10k	gender	half	official	bib	country	overall
15742	133.000	50	Moses, Julie K.	401	51.700	F	111.280	239.480	19713		
179	92.480	22	O'Connell, Thomas P II	2343	35.930	M	76.880	192.580	205		
11432	123.830	26	Venturelli, Gina M.	2226	49.620	F	104.120	218.080	13919		
16157	141.080	61	Cortes, Dora	25	57.820	F	118.970	237.480	20834		
895	98.570	50	Hrynowski, E. J.	12	39.350	M	83.220	167.870	1119		

```
'dataframe.dtypes\n'
```

```

25k          float64
age          int64
name         object
division     int64
10k          float64
gender       object
half         float64
official     float64
bib          object
country      object
overall      int64
pace         float64
30k          float64
5k           float64
genderdiv    int64
20k          float64
35k          float64
city         object
40k          float64
gender_int   float64
bib_int      int64
5kpace       float64
10kpace      float64
20kpace      float64
halfpace     float64
25kpace      float64
30kpace      float64
35kpace      float64
40kpace      float64
officialpace float64
racestd      float64
raceavg      float64
sustainer    float64
age_scaled   float64
overall_scaled float64
pace_scaled  float64
official_scaled float64
dtype: object

```

```
'dataframe.describe()\n'
```

	25k	age	division	10k	half	official	overall	pace	30k
count	16164.000	16164.000	16164.000	16164.000	16164.000	16164.000	16164.000	16164.000	16164.000
mean	118.036	41.638	1100.967	46.655	99.133	208.159	8429.373	7.947	143.423
std	13.423	10.351	942.115	5.254	10.965	23.744	5052.024	0.906	16.344
min	0.000	18.000	1.000	0.000	0.000	85.530	1.000	3.270	0.000

25%	108.900	34.000	363.000	43.200	91.530	191.727	4061.750	7.320	132.170
50%	119.150	42.000	842.000	47.170	100.080	209.225	8247.500	7.980	144.700
75%	127.320	49.000	1560.000	50.280	106.850	225.230	12662.250	8.600	154.780
max	163.620	80.000	3834.000	66.680	138.670	284.230	17598.000	10.850	195.870

'Label: post data_demographics on X #2: Column names are:'

Index(['age_scaled', 'sustainer', 'gender_int', 'racestd'], dtype='object')


train_size = 0.9, X_train is 14547, and y_train is 14547

test_size = 0.1, X_test is 1617, and y_test is 1617

'Label: after training_test_set: columns of X_train going into affinity_propagation: : Column n

Index(['age_scaled', 'sustainer', 'gender_int', 'racestd'], dtype='object')

0.1 We are plotting the elbow method!



2019-07-25 23:22:57: In: run_it

Begin

2019-07-25 23:22:57: In: run_it We are running with a Classifier model

2019-07-25 23:22:57: In: run_kmeans

Begin

running with number of clusters = 3

```
'dataframe.isnull().sum()'
```

```
age_scaled      0
sustainer       0
gender_int      0
racestd         0
overall_scaled  0
cluster         0
dtype: int64
```

```
'dataframe.columns\n'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd', 'overall_scaled', 'cluster'], dtype=object)
```

```
'dataframe.head(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
6260	0.565	0.000	1.000	0.129	0.458	1
10551	0.613	0.000	1.000	0.049	0.557	0
2327	0.097	0.000	1.000	0.309	0.656	0
13087	0.177	0.000	1.000	0.126	0.937	1
8747	0.435	0.000	1.000	0.222	0.945	0
4313	0.500	0.000	1.000	0.028	0.181	0
5463	0.323	0.000	1.000	0.230	0.562	0
10232	0.468	0.000	0.000	0.119	0.585	0
7808	0.452	0.000	1.000	0.064	0.248	1
6497	0.371	0.000	1.000	0.126	0.705	1

```
'dataframe.sample(10)\n'
```


	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
3366	0.194	1.000	1.000	0.014	0.113	0
12036	0.677	0.000	1.000	0.075	0.570	1
6540	0.355	0.000	1.000	0.151	0.348	2
7452	0.452	0.000	1.000	0.101	0.450	2
2238	0.274	0.000	1.000	0.113	0.151	1
2635	0.226	1.000	1.000	0.012	0.071	1
2972	0.484	0.000	1.000	0.138	0.602	0
14544	0.484	0.000	0.000	0.102	0.560	1
2719	0.419	0.000	1.000	0.124	0.243	0
909	0.145	0.000	1.000	0.211	0.135	1

```
'dataframe.dtypes\n'
```

```
age_scaled      float64
sustainer        float64
gender_int       float64
racestd          float64
overall_scaled   float64
cluster          int32
dtype: object
```

```
'dataframe.describe()\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
count	13107.000	13107.000	13107.000	13107.000	13107.000	13107.000
mean	0.364	0.057	0.645	0.123	0.439	0.671
std	0.162	0.232	0.478	0.165	0.273	0.573
min	0.000	0.000	0.000	0.003	0.000	0.000
25%	0.242	0.000	0.000	0.045	0.208	0.000
50%	0.371	0.000	1.000	0.083	0.419	1.000
75%	0.484	0.000	1.000	0.149	0.648	1.000
max	0.919	1.000	1.000	2.756	1.000	2.000



'next plot please'

'Label: columns of X_train going into affinity_propagation: : Column names are:'

Index(['age_scaled', 'sustainer', 'gender_int', 'racestd'], dtype='object')

2019-07-25 23:22:58: In: run_affinity_propagation
starting AffinityPropagation

2019-07-25 23:22:58: In: run_affinity_propagation
Begin

```
Estimated number of clusters: 7625
from run_affinity_propagation 0.11268298691068244
```

```
2019-07-25 23:35:29: In: run_affinity_propagation
finished with AffinityPropagation
```

```
'dataframe.isnull().sum()'
```

```
age_scaled      0
sustainer       0
gender_int      0
racestd         0
overall_scaled  0
cluster         0
dtype: int64
```

```
'dataframe.columns\n'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd', 'overall_scaled', 'cluster'], dtype=object)
```

```
'dataframe.head(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
6260	0.565	0.000	1.000	0.129	0.458	5421
10551	0.613	0.000	1.000	0.049	0.557	5601
2327	0.097	0.000	1.000	0.309	0.656	1242
13087	0.177	0.000	1.000	0.126	0.937	6884
8747	0.435	0.000	1.000	0.222	0.945	4624
4313	0.500	0.000	1.000	0.028	0.181	2285
5463	0.323	0.000	1.000	0.230	0.562	2921
10232	0.468	0.000	0.000	0.119	0.585	5433
7808	0.452	0.000	1.000	0.064	0.248	4138
6497	0.371	0.000	1.000	0.126	0.705	1431

```
'dataframe.sample(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
13848	0.290	0.000	0.000	0.033	0.537	7267
7118	0.645	0.000	1.000	0.239	0.492	187
6082	0.419	0.000	1.000	0.074	0.364	3233
14123	0.597	0.000	1.000	0.171	0.754	7409
7073	0.516	0.000	1.000	0.058	0.331	4977

9789	0.355	0.000	0.000	0.116	0.533	4734
7421	0.065	0.000	1.000	0.046	0.378	3949
3871	0.210	0.000	1.000	0.063	0.175	2055
770	0.129	0.000	1.000	0.039	0.015	3757
1299	0.242	1.000	1.000	0.011	0.076	5922

'dataframe.dtypes\n'


```

age_scaled      float64
sustainer       float64
gender_int      float64
racestd         float64
overall_scaled  float64
cluster         int64
dtype: object

```

'dataframe.describe()\n'

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
count	13107.000	13107.000	13107.000	13107.000	13107.000	13107.000
mean	0.364	0.057	0.645	0.123	0.439	3605.407
std	0.162	0.232	0.478	0.165	0.273	2171.660
min	0.000	0.000	0.000	0.003	0.000	0.000
25%	0.242	0.000	0.000	0.045	0.208	1698.500
50%	0.371	0.000	1.000	0.083	0.419	3732.000
75%	0.484	0.000	1.000	0.149	0.648	5296.500
max	0.919	1.000	1.000	2.756	1.000	7624.000



2019-07-25 23:35:30: In: run_mean_shift
Begin

from mean shift 0.7359724125900627
Number of estimated clusters: 11

2019-07-25 23:35:33: In: run_mean_shift
End

'dataframe.isnull().sum()'

age_scaled	0
sustainer	0
gender_int	0
racestd	0

```
overall_scaled    0
cluster           0
dtype: int64
```

```
'dataframe.columns\n'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd', 'overall_scaled', 'cluster'], dtype=object)
```

```
'dataframe.head(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
6260	0.565	0.000	1.000	0.129	0.458	0
10551	0.613	0.000	1.000	0.049	0.557	1
2327	0.097	0.000	1.000	0.309	0.656	1
13087	0.177	0.000	1.000	0.126	0.937	0
8747	0.435	0.000	1.000	0.222	0.945	1
4313	0.500	0.000	1.000	0.028	0.181	1
5463	0.323	0.000	1.000	0.230	0.562	1
10232	0.468	0.000	0.000	0.119	0.585	1
7808	0.452	0.000	1.000	0.064	0.248	0
6497	0.371	0.000	1.000	0.126	0.705	0

```
'dataframe.sample(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
12243	0.210	0.000	0.000	0.162	0.720	0
9601	0.629	0.000	1.000	0.052	0.388	2
1328	0.419	0.000	1.000	0.031	0.063	1
13866	0.468	0.000	0.000	0.082	0.654	4
8247	0.548	0.000	1.000	0.025	0.271	1
7378	0.516	0.000	1.000	0.044	0.309	1
7635	0.516	0.000	1.000	0.100	0.407	0
12871	0.097	0.000	0.000	0.056	0.570	0
99	0.161	0.000	1.000	0.070	0.002	0
4305	0.306	0.000	1.000	0.043	0.250	0


```
'dataframe.dtypes\n'
```

```
age_scaled      float64
sustainer        float64
gender_int       float64
```

```
racestd          float64
overall_scaled   float64
cluster          int64
dtype: object
```

```
'dataframe.describe()\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
count	13107.000	13107.000	13107.000	13107.000	13107.000	13107.000
mean	0.364	0.057	0.645	0.123	0.439	0.547
std	0.162	0.232	0.478	0.165	0.273	0.782
min	0.000	0.000	0.000	0.003	0.000	0.000
25%	0.242	0.000	0.000	0.045	0.208	0.000
50%	0.371	0.000	1.000	0.083	0.419	0.000
75%	0.484	0.000	1.000	0.149	0.648	1.000
max	0.919	1.000	1.000	2.756	1.000	10.000



'Label: entering run_spectral_clustering, data has shape of:: Dataframe shape is:'

(14547, 4)

'Label: entering run_spectral_clustering, target has shape of:: Dataframe shape is:'

(14547,)

2019-07-25 23:35:34: In: run_spectral_clustering

Begin

2019-07-25 23:35:34: In: run_spectral_clustering Running spectral_clustering with 3 clusters.

```
from spectral clustering 0.7531897945637207
```

```
2019-07-25 23:37:02: In: run_spectral_clustering  
End
```

```
'dataframe.isnull().sum()'
```

```
age_scaled      0  
sustainer       0  
gender_int      0  
racestd         0  
overall_scaled  0  
cluster         0  
dtype: int64
```

```
'dataframe.columns\n'
```

```
Index(['age_scaled', 'sustainer', 'gender_int', 'racestd', 'overall_scaled', 'cluster'], dtype=object)
```

```
'dataframe.head(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
6260	0.565	0.000	1.000	0.129	0.458	0
10551	0.613	0.000	1.000	0.049	0.557	2
2327	0.097	0.000	1.000	0.309	0.656	2
13087	0.177	0.000	1.000	0.126	0.937	0
8747	0.435	0.000	1.000	0.222	0.945	2
4313	0.500	0.000	1.000	0.028	0.181	2
5463	0.323	0.000	1.000	0.230	0.562	2
10232	0.468	0.000	0.000	0.119	0.585	2
7808	0.452	0.000	1.000	0.064	0.248	0
6497	0.371	0.000	1.000	0.126	0.705	0

```
'dataframe.sample(10)\n'
```

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
4175	0.597	0.000	1.000	0.204	0.523	2
4671	0.290	0.000	1.000	0.087	0.208	0
14179	0.419	0.000	0.000	0.042	0.678	2
8402	0.484	0.000	0.000	0.392	0.918	0
14158	0.613	0.000	1.000	0.249	0.995	2
8751	0.500	1.000	1.000	0.017	0.208	0

1938	0.468	0.000	1.000	0.024	0.206	2
11113	0.468	0.000	0.000	0.200	0.662	0
11380	0.210	0.000	0.000	0.029	0.574	0
7436	0.323	0.000	0.000	0.096	0.378	2

'dataframe.dtypes\n'


```

age_scaled      float64
sustainer       float64
gender_int      float64
racestd         float64
overall_scaled  float64
cluster         int32
dtype: object

```

'dataframe.describe()\n'

	age_scaled	sustainer	gender_int	racestd	overall_scaled	cluster
count	13107.000	13107.000	13107.000	13107.000	13107.000	13107.000
mean	0.364	0.057	0.645	0.123	0.439	0.820
std	0.162	0.232	0.478	0.165	0.273	0.956
min	0.000	0.000	0.000	0.003	0.000	0.000
25%	0.242	0.000	0.000	0.045	0.208	0.000
50%	0.371	0.000	1.000	0.083	0.419	0.000
75%	0.484	0.000	1.000	0.149	0.648	2.000
max	0.919	1.000	1.000	2.756	1.000	2.000



2019-07-25 23:37:02: In: run_it End
2019-07-25 23:37:02: In: main Ending main()

[]: